AQUAQ Analytics

Optane

Experts in fast data solutions

for demanding environments

AQUAQ Analytics

Optane AppDirect support

What is Optane?

- New type of RAM –NVRAM (Non-Volatile) or persistent RAM
- Unlike DRAM, contents persist after machine restart – more like disk
- Slower than DRAM, much faster than SSD
- Cheaper than DRAM and maximum size per chip much higher

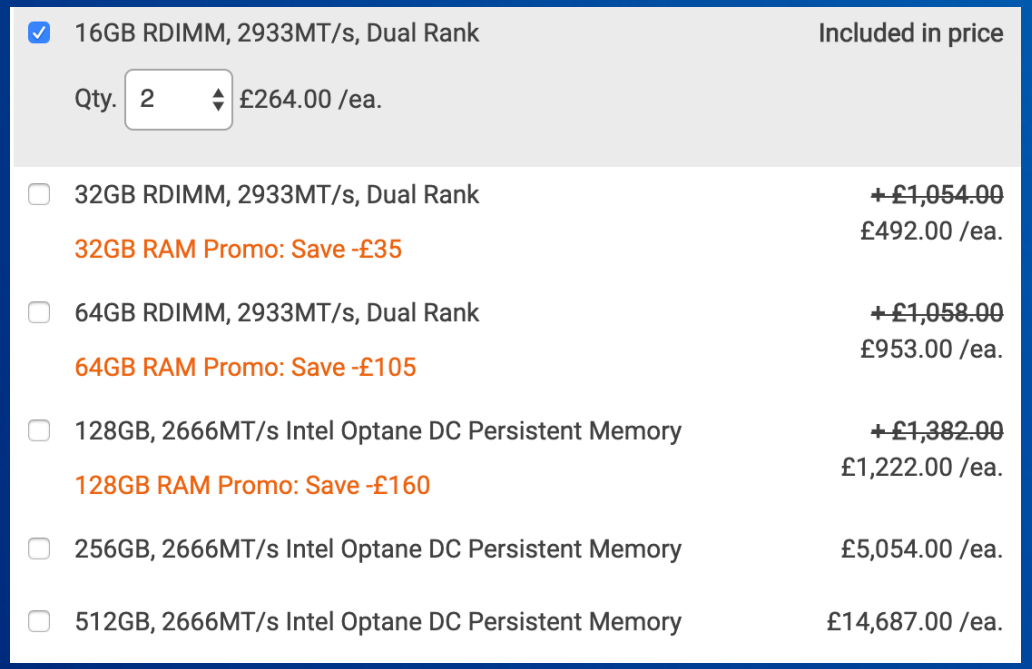

AQUAQ Analytics

Optane AppDirect support

Optane Performance

DDR4 memory accesses – 14ns
Optane DIMM – 350ns
NVMe Optane SSD access can take 10,000ns (10µs)
NVMe NAND SSD write – 30,000ns (30µs)
NVMe NAND SSD read – 120,000ns (120µs)
SATA NAND SSD read – 500,000ns (500µs or 0.5ms)
SATA NAND SSD write – 3,000,000ns (3,000µs or 3ms)
Disk drive seek – 100,000,000ns (100,000µs or 100ms)

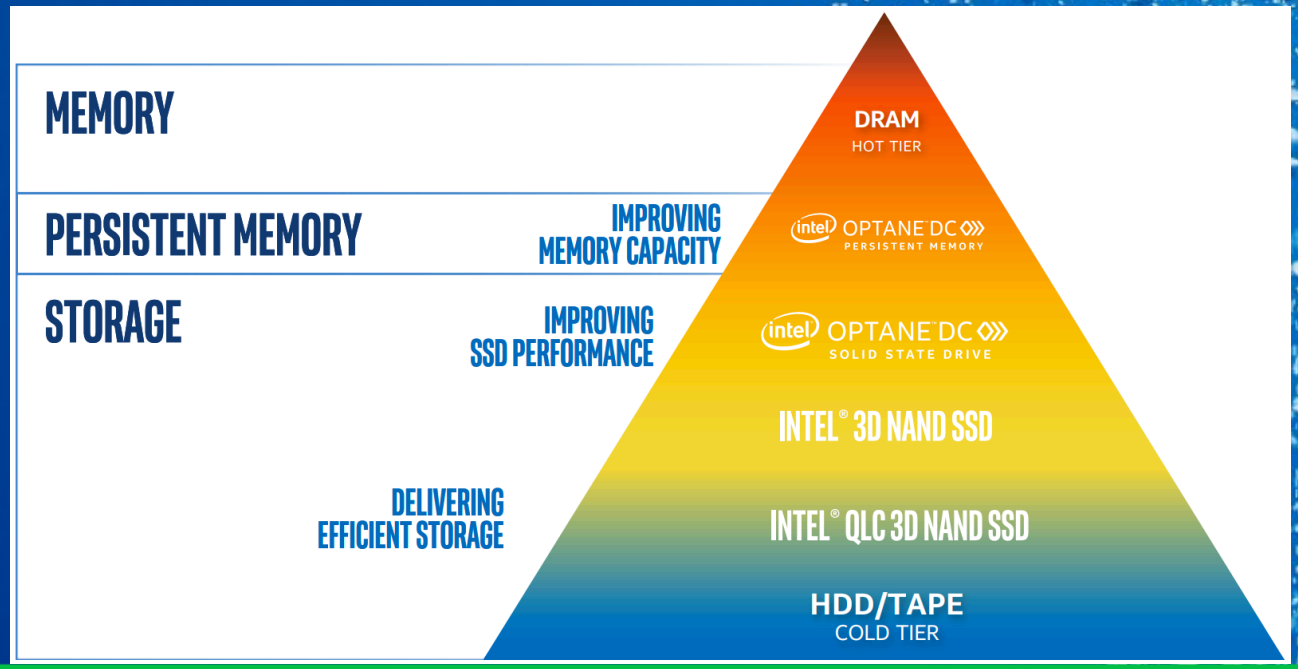

AQUAQ Analytics

Optane AppDirect support

Optane Modes

- Storage mode – Optane presents as a disk
- Memory mode – Optane becomes main memory pool, DRAM is L4 cache
- AppDirect mode – DRAM and Optane present as separate memory pools to applications

AQUAQ Analytics

Optane AppDirect support

Storage Mode

- Applications talk to Optane through a file system interface
- Optane just behaves as very fast SSD
- File API impacts performance – not as fast as direct access
- Transparent to kdb+, compatible with all versions

AQUAQ Analytics

Optane AppDirect support

Memory Mode

- Optane becomes the main memory pool – so a machine with 256Gb of DRAM and 1Tb of Optane would appear to applications as having 1Tb RAM
- DRAM is managed as cache transparently by OS
- Frequently accessed objects in memory should remain in DRAM with occasional cache misses which have to hit Optane
- Compatible with all versions of kdb+ - could potentially have huge 1Tb+ RDBs if caching performance is suited to how kdb+ is accessing in memory data

AQUAQ Analytics

Optane AppDirect support

App Direct Mode

- Applications can see DRAM and Optane pools separately
- Software must be rewritten to take advantage of Optane memory
- Kdb+ 4.0 adds support for App Direct mode with compatible Intel processor (Cascade Lake - 2020)

AQUAQ Analytics

Benchmarks - setup

- Intel loaner machine – 48 core, 384Gb DRAM, 6Tb Optane
- Generated quotes and trades database
- 4096 syms (`aaa, `aab, … , `ppp)
- 1.3B quotes, 300M trades per day, 5 days
- On disk: 70G per day, 350G total

AQUAQ Analytics

Benchmarks - setup

```
[user1@atsnode24 ~]$ sudo ipmctl show -topology
[sudo] password for user1:
 DimmID | MemoryType                  | Capacity    | PhysicalID| DeviceLocator
==============================================================================
 0x0001 | Logical Non-Volatile Device | 502.563 GiB | 0x0026    | CPU1_DIMM_A2
 0x0011 | Logical Non-Volatile Device | 502.563 GiB | 0x0028    | CPU1_DIMM_B2
 0x0021 | Logical Non-Volatile Device | 502.563 GiB | 0x002a    | CPU1_DIMM_C2
 0x0101 | Logical Non-Volatile Device | 502.563 GiB | 0x002c    | CPU1_DIMM_D2
 0x0111 | Logical Non-Volatile Device | 502.563 GiB | 0x002e    | CPU1_DIMM_E2
 0x0121 | Logical Non-Volatile Device | 502.563 GiB | 0x0030    | CPU1_DIMM_F2
 0x1001 | Logical Non-Volatile Device | 502.563 GiB | 0x0032    | CPU2_DIMM_A2
 0x1011 | Logical Non-Volatile Device | 502.563 GiB | 0x0034    | CPU2_DIMM_B2
 0x1021 | Logical Non-Volatile Device | 502.563 GiB | 0x0036    | CPU2_DIMM_C2
 0x1101 | Logical Non-Volatile Device | 502.563 GiB | 0x0038    | CPU2_DIMM_D2
 0x1111 | Logical Non-Volatile Device | 502.563 GiB | 0x003a    | CPU2_DIMM_E2
 0x1121 | Logical Non-Volatile Device | 502.563 GiB | 0x003c    | CPU2_DIMM_F2
 N/A    | DDR4                        | 32.000 GiB  | 0x0025    | CPU1_DIMM_A1
 N/A    | DDR4                        | 32.000 GiB  | 0x0027    | CPU1_DIMM_B1
 N/A    | DDR4                        | 32.000 GiB  | 0x0029    | CPU1_DIMM_C1
 N/A    | DDR4                        | 32.000 GiB  | 0x002b    | CPU1_DIMM_D1
 N/A    | DDR4                        | 32.000 GiB  | 0x002d    | CPU1_DIMM_E1
 N/A    | DDR4                        | 32.000 GiB  | 0x002f    | CPU1_DIMM_F1
 N/A    | DDR4                        | 32.000 GiB  | 0x0031    | CPU2_DIMM_A1
 N/A    | DDR4                        | 32.000 GiB  | 0x0033    | CPU2_DIMM_B1
 N/A    | DDR4                        | 32.000 GiB  | 0x0035    | CPU2_DIMM_C1
 N/A    | DDR4                        | 32.000 GiB  | 0x0037    | CPU2_DIMM_D1
 N/A    | DDR4                        | 32.000 GiB  | 0x0039    | CPU2_DIMM_E1
 N/A    | DDR4                        | 32.000 GiB  | 0x003b    | CPU2_DIMM_F1
[user1@atsnode24 ~]$
[user1@atsnode24 ~]$
[user1@atsnode24 ~]$ sudo ipmctl show -memoryresources
 MemoryType   | DDR         | PMemModule    | Total
============================================================
 Volatile     | 381.500 GiB | 0.000 GiB     | 381.500 GiB
 AppDirect    | -           | 6024.000 GiB  | 6024.000 GiB
 Cache        | 0.000 GiB   | -             | 0.000 GiB
 Inaccessible | 2.500 GiB   | 7.184 GiB     | 9.684 GiB
 Physical     | 384.000 GiB | 6031.184 GiB  | 6415.184 GiB
```

AQUAQ Analytics

Benchmarks - setup

```
ipmctl create -goal persistentmemorytype=appdirect

ndctl create-namespace --mode=fsdax --region=0 --size=2052G --align=2M

mkfs -t xfs /dev/pmem0

mkdir /mnt/pmem
mount -o dax /dev/pmem0 /mnt/pmem/
chmod 777 /mnt/pmem
```

AQUAQ Analytics

Benchmarks – Optane vs DRAM

```
/ Queries

/ simple select
\t:100 select from trades where sym in -5?`3
/ simple agg
\t:100 select avg price, sum size by sym from trades where sym in -10?`3
/ asof join
\t:10 aj[`sym`time;select from trades where sym in -5?`3;quotes]
/ snapshot
\t:10 select by sym from trades where sym in -100?`3,time<=2014.04.21D10
```

AQUAQ Analytics

Benchmarks – Optane vs DRAM

Scenario 1

- Single kdb+ process:
  - 1 day of data in DRAM
  - 1 day of data in Optane

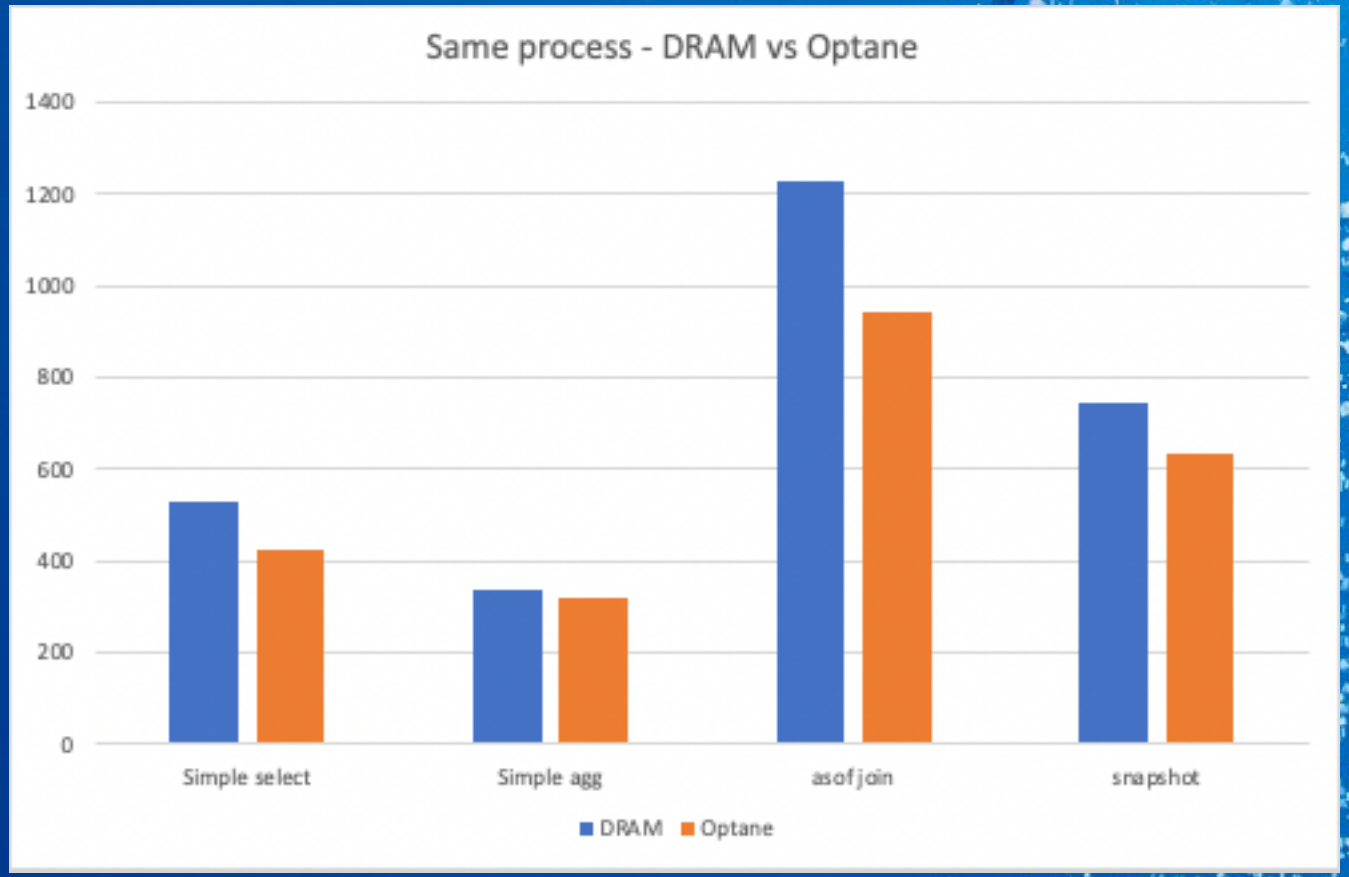

AQUAQ Analytics

Benchmarks – Optane vs DRAM

Scenario 2

- Separate kdb+ processes:
  - 1 day of data in DRAM
  - 1 day of data in Optane

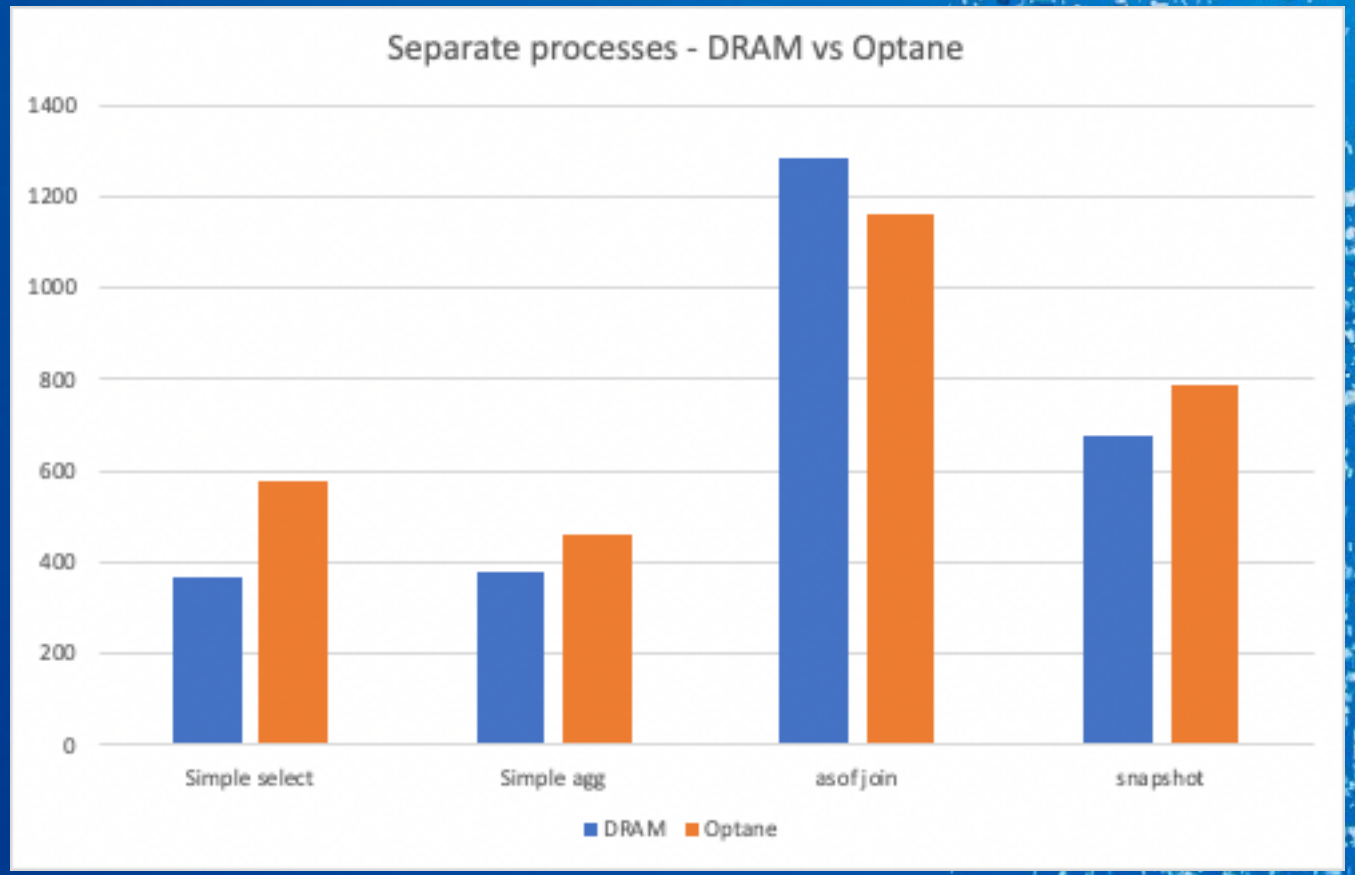

AQUAQ Analytics

Benchmarks – Optane vs Disk

- 5 days history loaded in Optane
- Each table is nested by date:

```
q)trades
2014.04.21| +`sym`time`src`price`size!(`p#`aaa`aaa`aaa`aaa`aaa`aaa`aaa`aaa`aa..
2014.04.22| +`sym`time`src`price`size!(`p#`aaa`aaa`aaa`aaa`aaa`aaa`aaa`aaa`aa..
2014.04.23| +`sym`time`src`price`size!(`p#`aaa`aaa`aaa`aaa`aaa`aaa`aaa`aaa`aa..
2014.04.24| +`sym`time`src`price`size!(`p#`aaa`aaa`aaa`aaa`aaa`aaa`aaa`aaa`aa..
2014.04.25| +`sym`time`src`price`size!(`p#`aaa`aaa`aaa`aaa`aaa`aaa`aaa`aaa`aa..
q){d:.Q.w[];system"d .m";o:.Q.w[];system"d .";(`used`mphy#(d;o))%1024 xexp 3}[]
used         mphy
---------------------
0.0003374666 375.1356
552          2018.949
```

AQUAQ Analytics

Benchmarks – Optane vs Disk

```
/ Queries

/ simple select - 3 days, 10 syms
\t select from quotes where date in -3?date, sym in -10?`3
\t raze {select from quotes[x] where sym in -10?`3} each -3?key quotes

/ aggregation - 3 days, 5 syms, 1h bars
\t select last bid, last ask by date, sym, 0D1 xbar time from quotes where date in -3?date, sym in -5?`3
\t raze {select last bid, last ask by sym, 0D1 xbar time from quotes[x] where sym in -5?`3} each -3?key quotes

/ asof lookup - 1 day, rack of 3000 sym & time samples
rack:([]time:2014.04.23D09:30+0D00:05*til 60)cross([]sym:-50?`3)
\t:50 aj[`sym`time;rack;select time,sym,bid,ask from quotes where date=2014.04.23]
\t:50 aj[`sym`time;rack;quotes[2014.04.23]]
```

AQUAQ Analytics

Benchmarks – Optane vs Disk

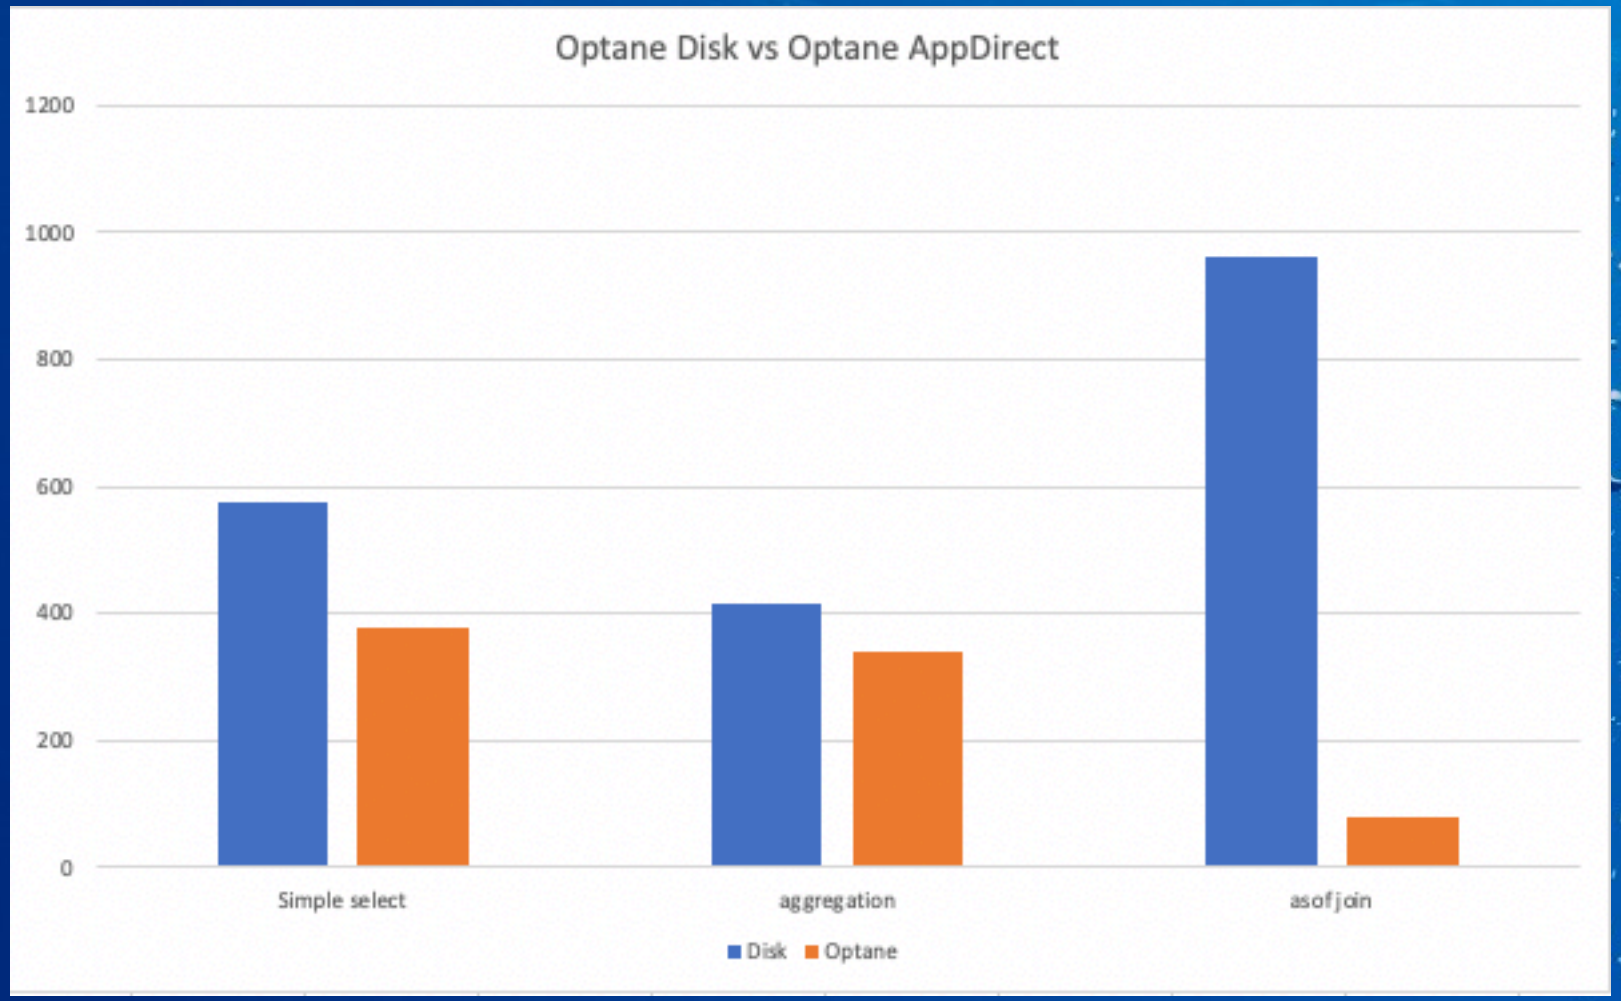

AQUAQ Analytics

Conclusions

- Performance is great – Optane Appdirect delivers DRAM-like performance for most kdb+ workloads tested
- Largest Optane chips are expensive, but deliver memory in a volume not possible with DRAM
- Flexible – can be chopped up into partitions to separate process memory pools and used as disk at the same time
- Kdb+ architectures will have to change to take advantage of Optane

AQUAQ Analytics

Q&A