



AquaQuarantine: TorQ Tickerplant

Experts in fast data solutions
for demanding environments

- Established in 2011
- Headquarters in Belfast, N.Ireland
- Headcount of 160 staff
- 2016 US Subsidiary launched
- 2018 Singapore subsidiary launch
- 2020 Hong Kong subsidiary launch



What do we do?

Technology
Consultancy Services




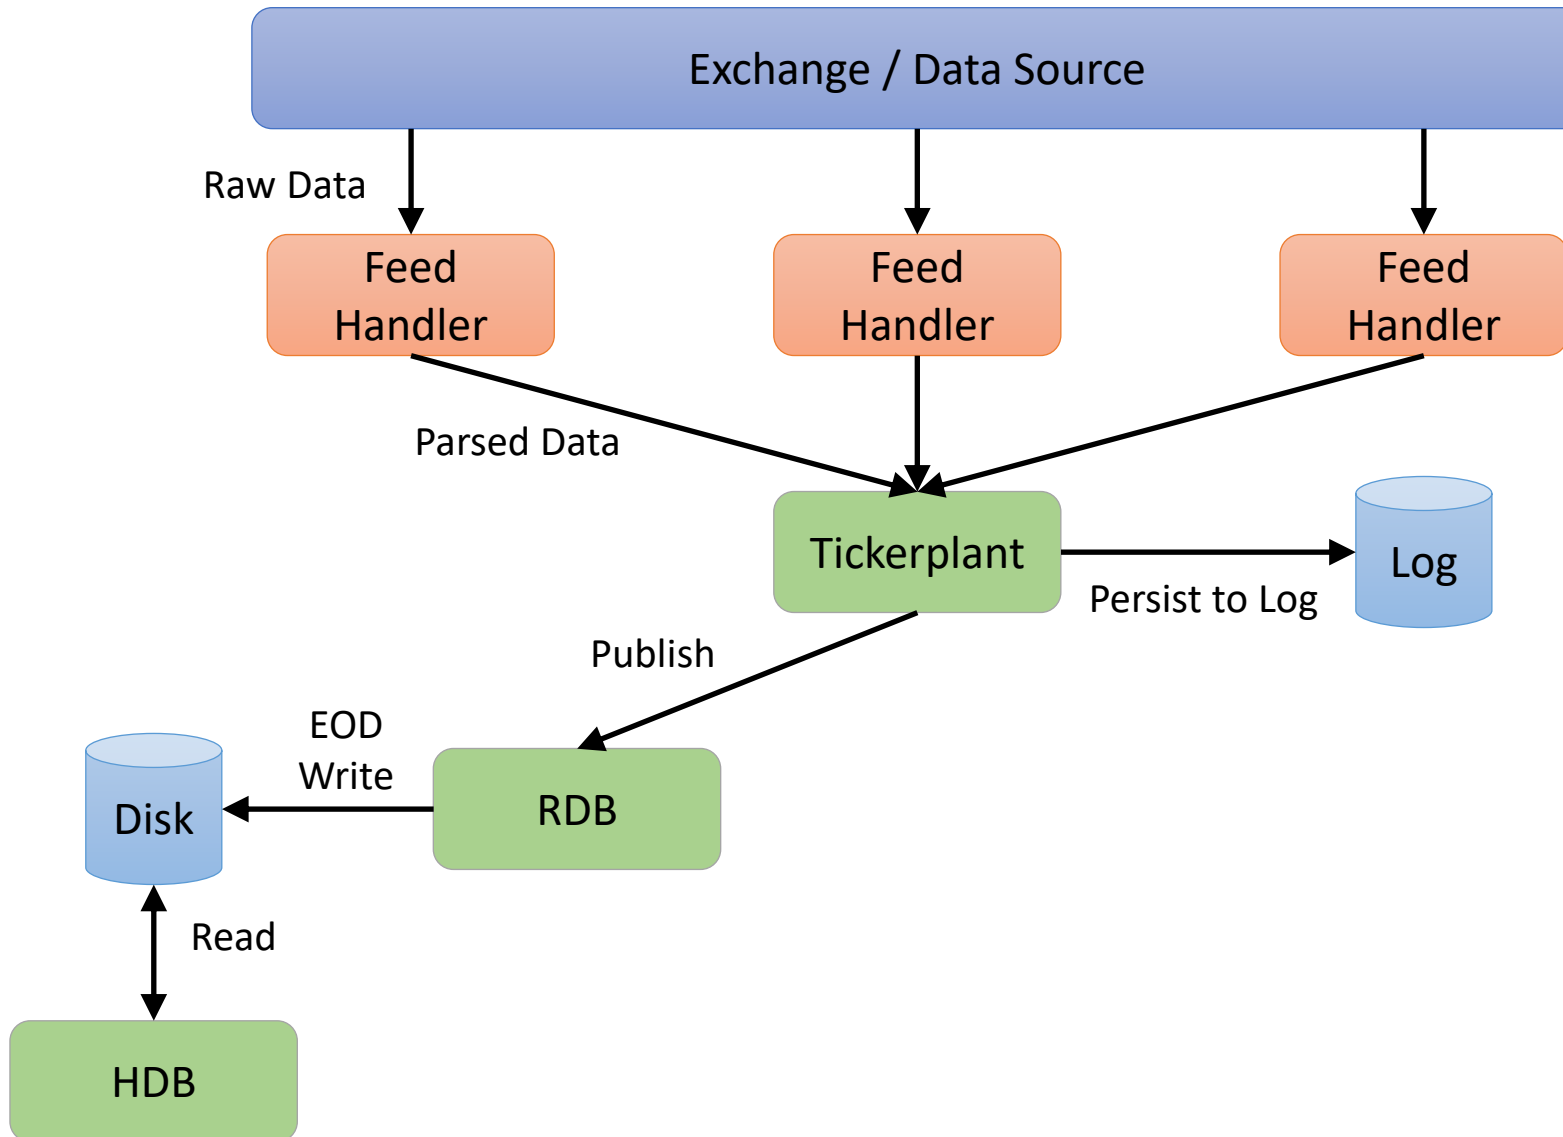
Altair Panopticon
Professional Services



Remote (24/7) Support
Centre of Excellence




Why?



Why?

- TorQ Tickerplant is the kdb+tick TP with some minor modifications
- In various installations we've tweaked some things, or thought "wouldn't it be nice if..."
- Main drivers are:
 - Greater flexibility
 - Easier customisation
 - Easier to support
 - Building for future enhancements




Backward Compatibility


- Need to ensure backward compatibility is maintained
- Maintain interface:
 - .u.upd/.u.sub/.u.add/log file format
- TorQ TP allows configuration of:
 - Immediate or Batch publish
 - Data Timezone
 - Roll Timezone + Roll Time Offset e.g. FX: timestamp in UTC, roll at 5pm NYC local




1. More granular log files



- Current TP has one log file per date
- Pros:
 - Simple
 - Message sequence is maintained
- Cons:
 - Can become unwieldy
 - Whole file must be replayed for subset of data
 - On restart log must be counted



- Proposed solution is to provide multiple logging formats
 - `periodic- single log for all tables, rolled on a configurable period
 - `tabular- single log for per table, rolls daily
 - `tabperiod- single log per table per period
 - `none- single log, rolls daily
 - `custom- different logging method per table, or no logging for some tables
 - Easy to extend e.g. per instrument log files (might be useful in FX)
- In all cases, we have a directory of log files



Default (tabperiod)

```

stplogs/
├── database2020.07.23/
│   ├── trade20200723090000
│   ├── trade20200723100000
│   ├── quote20200723090000
│   └── quote20200723100000
├── database2020.07.24/
│   ├── trade20200724090000
│   ├── trade20200724100000
│   ├── quote20200724090000
│   └── quote20200724100000

```

Periodic

```

stplogs/
├── database2020.07.23/
│   ├── periodic20200723090000
│   ├── periodic20200723100000
│   ├── periodic20200723110000
│   └── periodic20200723120000
├── database2020.07.24/
│   ├── periodic20200724090000
│   ├── periodic20200724100000
│   ├── periodic20200724110000
│   └── periodic20200724120000

```

Tabular

```

stplogs/
├── database2020.07.23/
│   ├── trade20200723000000
│   └── quote20200723000000
├── database2020.07.24/
│   ├── trade20200724000000
│   └── quote20200724000000


```

Lognames timestamped with start of logging period


- We also maintain a meta file at the top level of directory containing meta info on each file. Log file directory is self describing
- Meta info is:
 - Time period the log file covers
 - Tables contained in log file
 - Message Count (on best efforts basis)
 - Schema of tables contained

```
(q)last .stpm.metatable
seq      | 18i
logname  | `:/home/tp/test/stplogs/database2020.08.10/trade_iex20200810140000
start    | 2020.08.10D14:00:00.690988000
end      | 0Np
tbls     | ,`trade_iex
msgcount | 0i
schema   | (,`trade_iex)!,+`time`sym`price`size`stop`cond`ex`srctime!(`timestamp$();`g#`symbol$();`float$();`int$();`boolean$();"";"";`timestamp$())
additional | ,()!()
```

- Clients must replay all the log files they need data from
- Schema info is good for:
 - Historic compatibility
 - Intraday schema changes
- If inter-table message sequencing is required, use `periodic or `none logging modes
- When the TP restarts it creates a new log file with 0 count (instant restart due to no replay)
- Subscribers have to handle both `.u.end and `.u.endp (end-of-period)



2. New Batch Publish Mode




- Support same options as current TP:
 - Immediate: log and publish immediately
 - Batch: log, publish on a timer
- Added `memorybatch` option
 - Log to memory
 - On timer, flush to log and publish
- Pros:
 - Better performance (less disk writes)
- Cons:
 - Additional failure case to manage (once it's received by TP it hasn't been logged)




- Easy to add own logging modes
- Need to add upd function and timer function

```
[q).stplg.upd
      | ()
memorybatch | {[t;x;now]
  t insert updtab[t] . (x;now);
}
defaultbatch| {[t;x;now]
  t insert x:.stplg.updtab[t] . (x;now);
  `..loghandles[t] enlist(`upd;t;x);
  // track tmp counts, and add these after publish
  @['.stplg.tmpmsgcount;t++;1];
  @['.stplg.tmprowcount;t++;count first x];
}
immediate   | {[t;x;now]
  x:updtab[t] . (x;now);
  `..loghandles[t] enlist(`upd;t;x);
  @['.stplg.msgcount;t++;1];
  @['.stplg.rowcount;t++;count first x];
  .stpps.pub[t;x]
}
```

```
[q).stplg.zts
      | ()
memorybatch | {
  {[t]
    if[count value t;
      `..loghandles[t] enlist (`upd;t;value flip value t);
      @['.stplg.msgcount;t++;1];
      @['.stplg.rowcount;t++;count value t];
      .stpps.pubclear[t]];
    }each .stpps.t;
  }
defaultbatch| {
  // publish and clear all tables, increment counts
  .stpps.pubclear[.stpps.t];
  // after data has been published, updated the counts
  .stplg.msgcount+:.stplg.tmpmsgcount;
  .stplg.rowcount+:.stplg.tmprowcount;
  // reset temp counts
  .stplg.tmpmsgcount:.stplg.tmprowcount:()!();
}
immediate   | {}
```



3. Flexible Update Handling



- Supply a dictionary of update functions per table
- Default definition is to add current timestamp (this is useful in most use cases)

```
[q).stplg.updtab
logmsg      | {(enlist(count first x)#y),x}
quote       | {(enlist(count first x)#y),x}
quote_iex   | {(enlist(count first x)#y),x}
trade       | {(enlist(count first x)#y),x}
trade_iex   | {(enlist(count first x)#y),x}
. =         | {(enlist(count first x)#y),x}
```

- Maintain sequence number which can be appended to incoming messages
- Allow publishers to send multiple updates in one message (which will get the same timestamp and sequence number)

```
[q).stp.upd
{[t;x]
  // snap the current time and check for period end
  .stplg.checkends now:.z.p+.eodtime.dailyadj;
  // Type check allows update messages to contain multiple tables/data
  $[0h<type t;
    .stplg.updmsg' [t;x;now];
    .stplg.updmsg[t;x;now]
  ];
  .stplg.seqnum+:1;
}
```

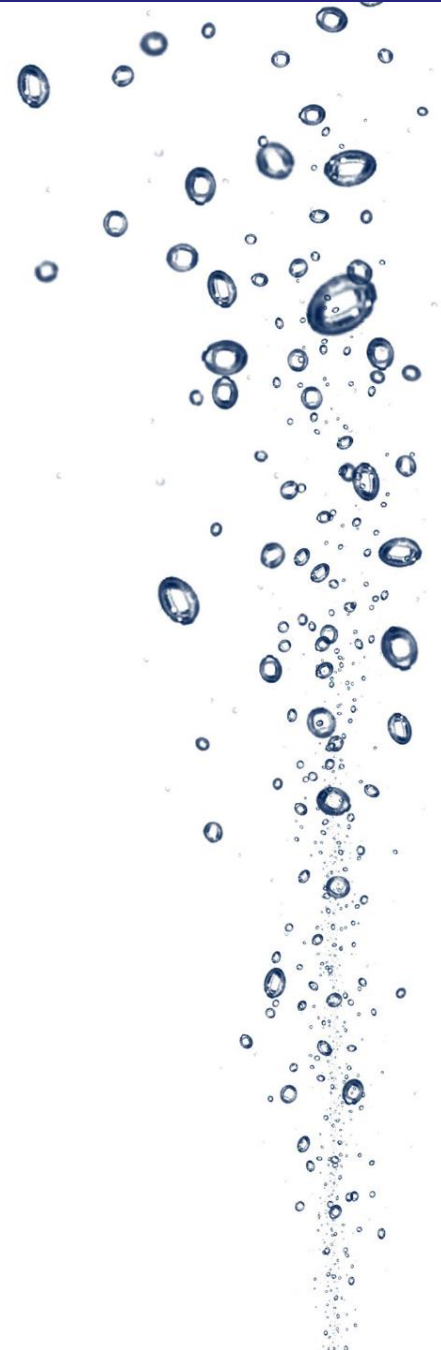
- Add error trap mode to capture failed messages to a separate error log
- Useful for debugging

```
[q).stplg.updtab  
logmsg    | {(enlist(count first x)#y),x}  
quote     | {(enlist(count first x)#y),x}  
quote_iex | {(enlist(count first x)#y),x}  
trade     | {(enlist(count first x)#y),x}  
trade_iex | {(enlist(count first x)#y),x}  
.-        | {(enlist(count first x)#y),x}
```


- Maintain both message count and row counts per table




4. Flexible Publishing




- Old TP pub/sub code was based on u.q from kdb+ tick, extended to incorporate broadcast publish (-25!)
- Clients could subscribe based on table and optionally sym
- New subscription functionality:
 - Supports same interface (.u.sub[tables;syms])
 - Utilises broadcast publish for subscriptions for all data from a table
 - Removes requirement for a table to have column called `sym
 - Allows any “where” filter to be applied
 - Allows columns to be specified
 - Result is less subscriber side filtering => lower overhead
- Complex where filtering should be reserved for non-primary capture tickerplants (e.g. chained tickerplants)




5. Compatibility



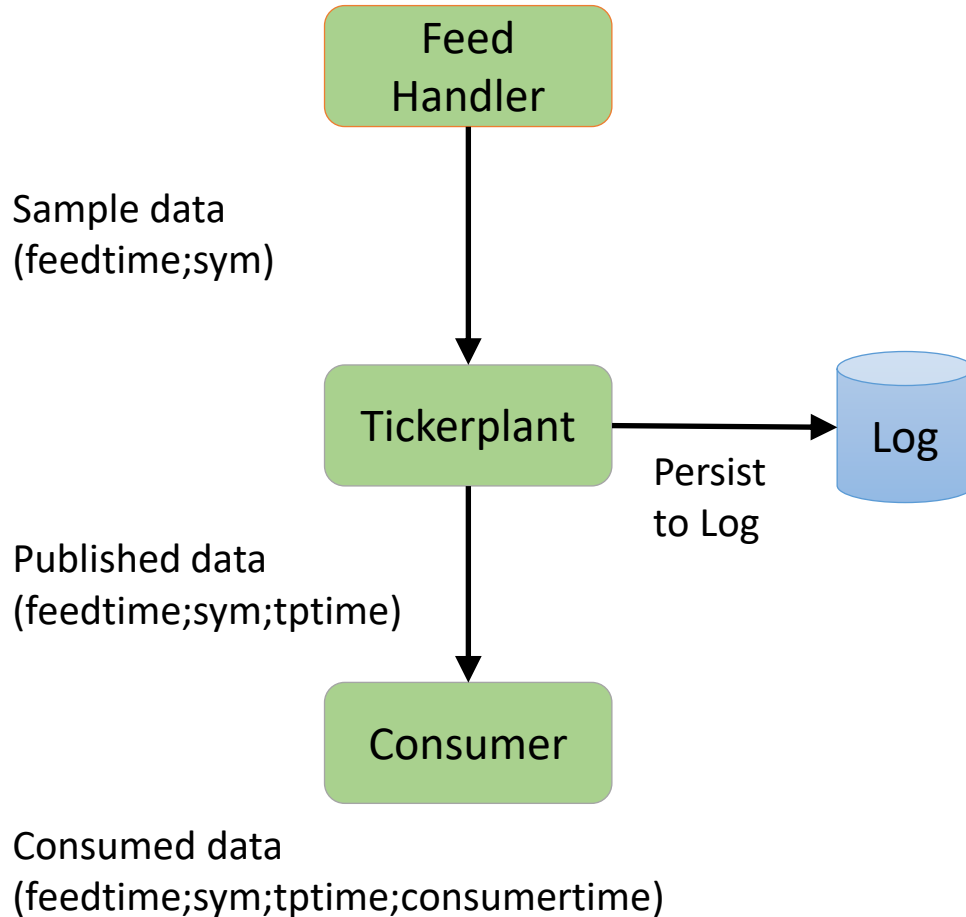
- All TorQ components will work with both old and new tickerplants
 - RDB
 - WDB
 - Chained TP
 - Subscribers
 - Tickerlog Replay
- Will automatically replay multiple log files
- Will utilise new subscription mechanism
- We would like to phase out use of the old tickerplant, but will leave in place



6. Performance



Test Harness



- Tests executed for:
 - Single inserts (1 record)
 - Bulk inserts (1000 records)
- Measurements taken for:
 - Messages per second
 - Latencies between each point
 - Drift
- Lots of variables in real life!
 - Disk performance
 - Message size
 - Number of consumers
 - Etc...




- Results below are for single inserts
- Bulk insert results were broadly similar (dependent on I/O)
- Error mode adds a small over head to every insert


Mode	Old TP (messages p/s)	Segmented TP
Immediate	60,000	66,000
Immediate+Error	N/A	65,000
Batch	127,000	100,000
Memory Batch	N/A	164,000

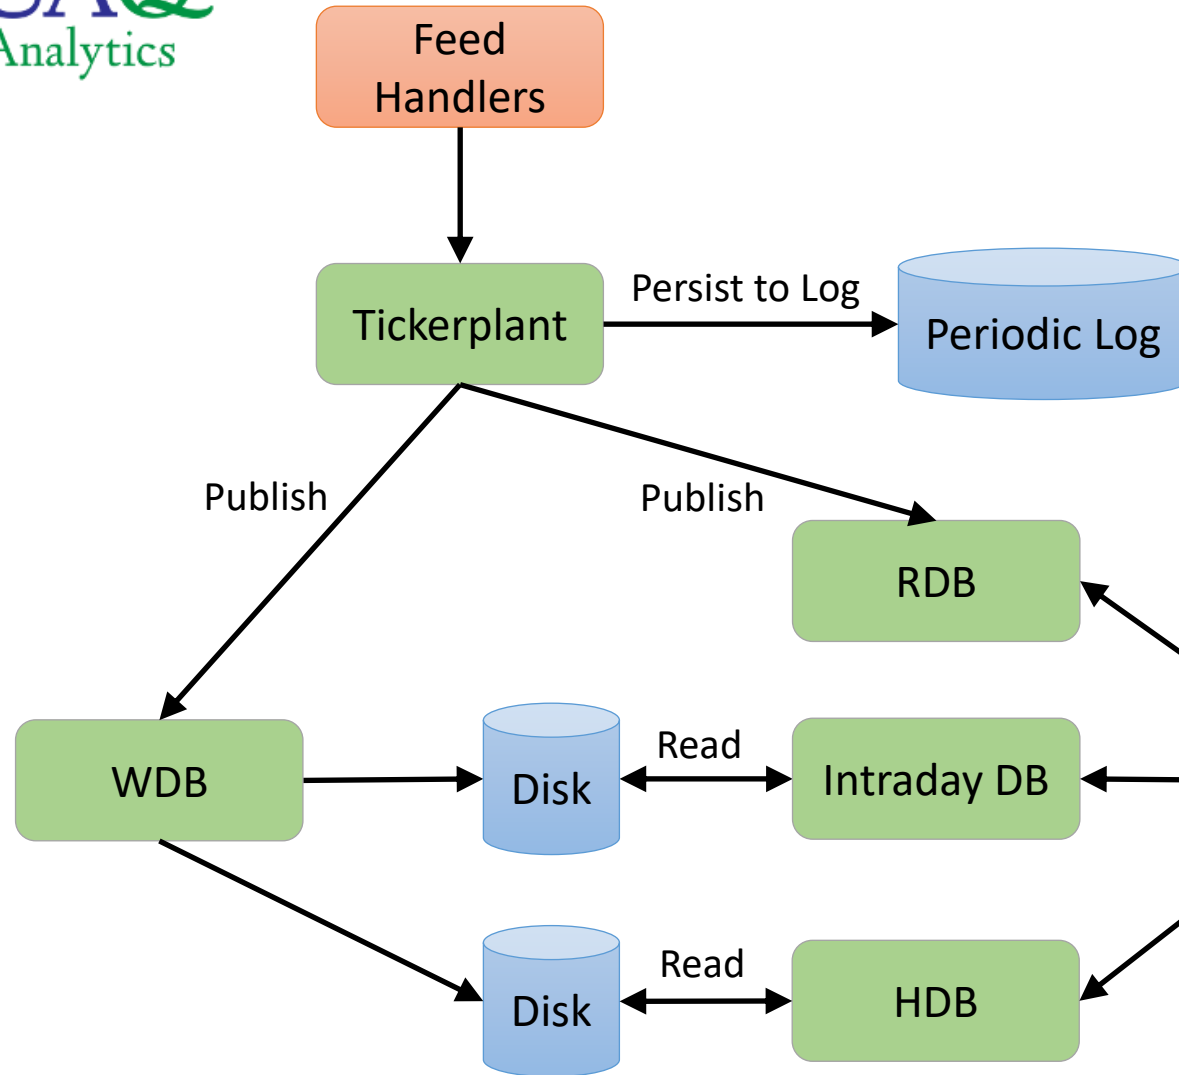


7. Future Planning



- Main driver was to allow more granular log files
 - Easier to manage
 - Targeted replays (table prioritisation)
- More frequent log files leads to additional design considerations
 - Old tickerplant had periodic log files where period = one day
 - One day aligns with date partitioning scheme of on disk database
 - More frequent log files lends itself to more granular writes
 - Also snapshots services become easier





- HDB = historic data, date partitioned
- Intraday DB = today's data prior to current period
- RDB = today's data since period start




- Intraday DB format:
 - Must be easy to append new data intraday
 - Should be well structured for efficiency of majority of queries
- Int partitioned to represent period YYYYMMDDHHMM e.g. 202008071400
- Int partitions mapped via a table to timeperiod+other set of identifiers

q)intmap

periodstart	sym	int
2020.08.07D00:00:00.000000000	a	0
2020.08.07D00:00:00.000000000	b	1
2020.08.07D00:00:00.000000000	c	2
2020.08.07D01:00:00.000000000	a	3
2020.08.07D01:00:00.000000000	b	4
2020.08.07D01:00:00.000000000	c	5
2020.08.07D02:00:00.000000000	a	6
2020.08.07D02:00:00.000000000	b	7
2020.08.07D02:00:00.000000000	c	8

- API functions for data access across RDB/Intraday/HDB would be best

- If retention of HDB date partitioning format required then probably best to build that separately
- Could remove date partitioned HDB in some cases and just move to RDB and period partitioned HDB (continuous rolling HDB)
- Pros:
 - Lower overall memory usage
 - Faster recovery of RDB or other subscribers following failure
- Cons:
 - Greater complexity



Thanks!

Q+A

